

Graph CNN

A/Prof Richard Yi Da Xu

Yida.Xu@uts.edu.au

Wechat: aubedata


<https://github.com/roboticcam/machine-learning-notes>

University of Technology Sydney (UTS)

April 12, 2019

Given a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, the associated degree matrix D is a $n \times n$ diagonal matrix:

$$d_{i,j} := \deg(v_i) \text{ if } i = j, \quad d_{i,j} := 0 \text{ otherwise}$$



$$\begin{pmatrix} 4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

- ▶ $\deg(v_i)$ counts the number of times an edge terminates at that vertex
- ▶ in undirected graph, each new loop increases degree of vertex by **two**, see node 1 above

matrix representation of a graph

- ▶ Given a simple graph \mathcal{G} with n vertices, its Laplacian matrix $L^{n \times n}$:

$$L^{\mathcal{G}} = D - A$$

- ▶ D is the degree matrix and A is the adjacency matrix of the graph.

$$L_{i,j}^{\mathcal{G}} := \begin{cases} \deg(v_i) & \text{if } i = j \\ -1 & \text{if } i \neq j \text{ and } v_i \text{ is adjacent to } v_j \\ 0 & \text{otherwise} \end{cases}$$

- ▶ Since \mathcal{G} is a simple graph, A only contains 1s or 0s and its diagonal elements are all 0s.
- ▶ Symmetric normalized Laplacian:

$$L^{\text{sym}} := D^{-1/2} L D^{-1/2} = I - D^{-1/2} A D^{-1/2}$$

About Graph Laplacian

- ▶ For a twice differentiable function f on the euclidean space, the Laplacian of f , Δf is **div(grad(f))** (divergence of the gradient of f)
- ▶ consider Laplacian in two dimensions:

$$\Delta f(x, y) = \nabla \cdot \nabla f(x, y) = \frac{d^2 f(x, y)}{dx^2} + \frac{d^2 f(x, y)}{dy^2}$$

- ▶ knowing that in 1-d, second derivative is approximated by:

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - 2f(x) + f(x-h)}{h^2}$$

- ▶ then we can see the approximation of

$$\begin{aligned} \Delta f(x, y) &= \lim_{h \rightarrow 0} \frac{f(x+h, y) + f(x-h, y) + f(x, y+h) + f(x, y-h) - 4f(x, y)}{h^2} \\ &\approx \frac{f(x+h, y) + f(x-h, y) + f(x, y+h) + f(x, y-h) - 4f(x, y)}{h^2} \end{aligned}$$

- ▶ this means that $\Delta f(x, y)$ is a **local gradient averaging operator**, is zero at a smooth image

- ▶ Consider a function (or vector) \mathbf{f} over vertices \mathcal{V} , $\mathbf{f} : \mathcal{V} \rightarrow \mathbb{R}$. This function is indexed by vertices, say v_i is the node value for i^{th} node, i.e., \mathbf{f} is a vector, for example:

$$\mathbf{f} = \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix}$$

- ▶ then the product $\nabla \mathbf{f}$ is indexed by edges \mathcal{E} :
- ▶ each element of $\nabla \mathbf{f}$ is the difference between two end points of an edge e i.e. $v_i - v_j$:

$$\nabla \mathbf{f} = \begin{bmatrix} 1 & -1 & 0 \\ 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix} = \begin{bmatrix} v_1 - v_2 \\ v_2 - v_3 \end{bmatrix}$$

- ▶ Matrix ∇ acts like a difference or gradient operator on \mathcal{G}

- ▶ Consider a real function \mathbf{g} over edges \mathcal{E} , $\mathbf{g} : \mathcal{E} \rightarrow \mathbb{R}$ indexed by edges.
- ▶ then, the product $\nabla^T g^T$ is a vector indexed by \mathcal{V} .
- ▶ Value at i^{th} vertex:

$$(\mathbf{g}\nabla)_i = \sum_{e \text{ exit node } i} g_e - \sum_{e \text{ enters node } i} g_e$$

- ▶ $(\mathbf{g}\nabla)_i$ is the net outbound flow on the vertex v which is **divergence**

combine together

- ▶ Consider the quantity $\nabla^T \nabla \mathbf{f}$
- ▶ this is divergence of the gradient.
- ▶ the matrix $L^{\mathcal{G}} = \nabla^T \nabla$ is the **Graph Laplacian**:

$$\begin{aligned}\nabla^T \nabla &= \begin{bmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} 1 & -1 & 0 \\ 0 & 1 & -1 \end{bmatrix} = \begin{bmatrix} 1 & -1 & 0 \\ -1 & 2 & 1 \\ 0 & 1 & -1 \end{bmatrix} \\ \Rightarrow \nabla^T \nabla \mathbf{f} &= \begin{bmatrix} 1 & -1 & 0 \\ 1 & 2 & -1 \\ 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix} = \begin{bmatrix} v_1 - v_2 \\ 2v_2 + v_1 - v_3 \\ v_2 - v_3 \end{bmatrix}\end{aligned}$$

- ▶ for each i, j element of $L^{\mathcal{G}}$:

$$L_{i,j}^{\mathcal{G}} := \begin{cases} \deg(v_i) & \text{if } i = j \\ -1 & \text{if } i \neq j \text{ and } v_i \text{ is adjacent to } v_j \\ 0 & \text{otherwise} \end{cases}$$

- ▶ easy to see that $L^{\mathcal{G}} = D - A$:
- ▶ Laplacian is positive semidefinite:

$$\mathbf{f}^T \nabla^T \nabla \mathbf{f} = \|\nabla \mathbf{f}\|^2 = \sum_{(i,j) \in \mathcal{E}} (\mathbf{f}_i - \mathbf{f}_j)^2$$

- ▶ A Laplacian is written as:

$$L^{\mathcal{G}} \phi_k = \lambda_k \phi_k$$

- ▶ eigenvectors:

$$\langle \phi_k, \phi_{k'} \rangle = \delta_{k,k'}$$

- ▶ eigenvalues are non-negative:

$$0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$$

- ▶ eigendecomposition of graph Laplacian:

$$L^{\mathcal{G}} = \Phi^{\top} \Lambda \Phi$$

where:

$$\Phi = [\phi_1, \dots, \phi_n] \quad \Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$$

- **traditional** Fourier series of a function:

$$f(x) = \sum_{k \geq 0} \underbrace{\frac{1}{2\pi} \int_{-\pi}^{\pi} f(t) \exp^{-ikt} dt}_{\hat{f}_k = \langle \mathbf{f}, \exp^{-ikx} \rangle \text{ coefficient}} \underbrace{\exp^{-ikx}}_{\text{Fourier basis}}$$

- **graph** Fourier series: for the i^{th} element:

$$f_i = \sum_{k=1}^n \underbrace{\langle f, \phi_k \rangle}_{\hat{f}_k \text{ coefficient}} \underbrace{\phi_{k,i}}_{\text{graph Fourier basis}}$$

in matrix notation:

$$\hat{\mathbf{f}} = \Phi^{\top} \mathbf{f} \quad \text{and} \quad \mathbf{f} = \Phi^{\top} \hat{\mathbf{f}}$$

Convolution in Euclidean space

- convolution theorem:

$$\widehat{\mathbf{f} \circledast \mathbf{g}} = \hat{\mathbf{f}} \odot \hat{\mathbf{g}}$$

- $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\begin{aligned} (\mathbf{f} \circledast \mathbf{g})_i &= \sum_k g(i - k) \bmod n \cdot f_m \\ \Rightarrow \mathbf{f} \circledast \mathbf{g} &= \underbrace{\begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix}}_{\text{Circulant matrix}} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix} \\ &= \Phi \underbrace{\begin{bmatrix} \hat{g}_1 & & & \\ & \ddots & & \\ & & \ddots & \\ & & & \hat{g}_n \end{bmatrix}}_{\Phi^\top \mathbf{g}} \Phi^\top \mathbf{f} \\ &= \Phi(\Phi^\top \mathbf{g} \odot \Phi^\top \mathbf{f}) \end{aligned}$$

- ▶ spectral convolution:

$$\begin{aligned}\mathbf{f} \circledast \mathbf{g} &= \Phi(\Phi^\top \mathbf{g} \odot \Phi^\top \mathbf{f}) \\ &= \underbrace{\Phi \text{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_G \mathbf{f}\end{aligned}$$

What is missing?

- ▶ CNN has a pooling step
- ▶ GraphCNN has Graph downsampling(**coarsening**) step